

Outline

Module 1 : 大數據簡介

Module 2 : Hadoop Ecosystem介紹

Module 3 : Hadoop 平台安裝

Module 4 : Hadoop 分散式檔案系統 (HDFS)

Module 5 : Hadoop MapReduce

Module 6 : Apache Hive

Module 7 : Sqoop與Flume

Module 8 : Apache Spark

Module 9 : Spark 平台安裝

Module 10 : RDD – Resilient distributed dataset


Module 11 : Scala 程式開發基礎

Module 12 : Spark SQL 及 DataFrame

Module 13 : Spark 機器學習函式庫(MLlib)



MapReduce的由來

- ▶ 為處理大巨量資料而來－假設單一電腦最多可同時處理2百萬筆資料，要如何處理2億筆資料？
 - ~~分一百次處理~~
 - 分散至一百台電腦處理 

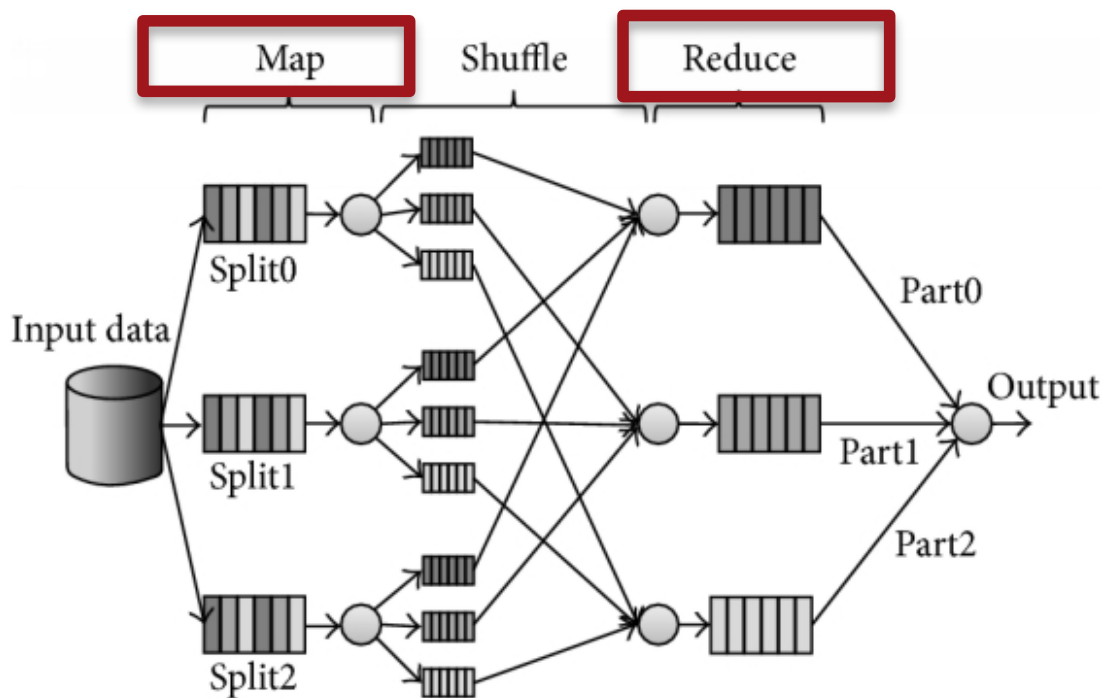
分而治之、各別擊破(Divide and Conquer)的概念

MapReduce的概念

- ▶ Google發表GFS及MapReduce相關論文(2003~2004)
- ▶ 處理巨量資料的共同課題：
 - 使用多台機器(cluster)共同處理
 - 要處理兩項基本作業：**Map**及**Reduce**
- ▶ Google主要受到**函數編程(Functional Programming)**中Map / Reduce的啟發
 - Map－將輸入**轉換為另一個集合**(方便後續處理)
 - Reduce－將Map的結果進行**聚合處理**(加總或計數等)
 - Ex：求[1,2,3,4,5]的平方和
 - Map： $[1,2,3,4,5] - (^2) \rightarrow [1,4,9,16,25]$
 - reduce： $[1,4,9,16,25] - (\text{sum}) \rightarrow 55$

MapReduce的軟體架構

- ▶ 由Google提出，在電腦叢集上執行分散式運算的軟體架構
- ▶ 開發人員只需專注於定義Map及Reduce的執行內容
 - Map必需實作，Reduce可視情況決定是否實作
- ▶ 平行運算的其它細節(如工作及資源分配、輸入的切分、輸出結果的收集等)由MapReduce框架負責協調



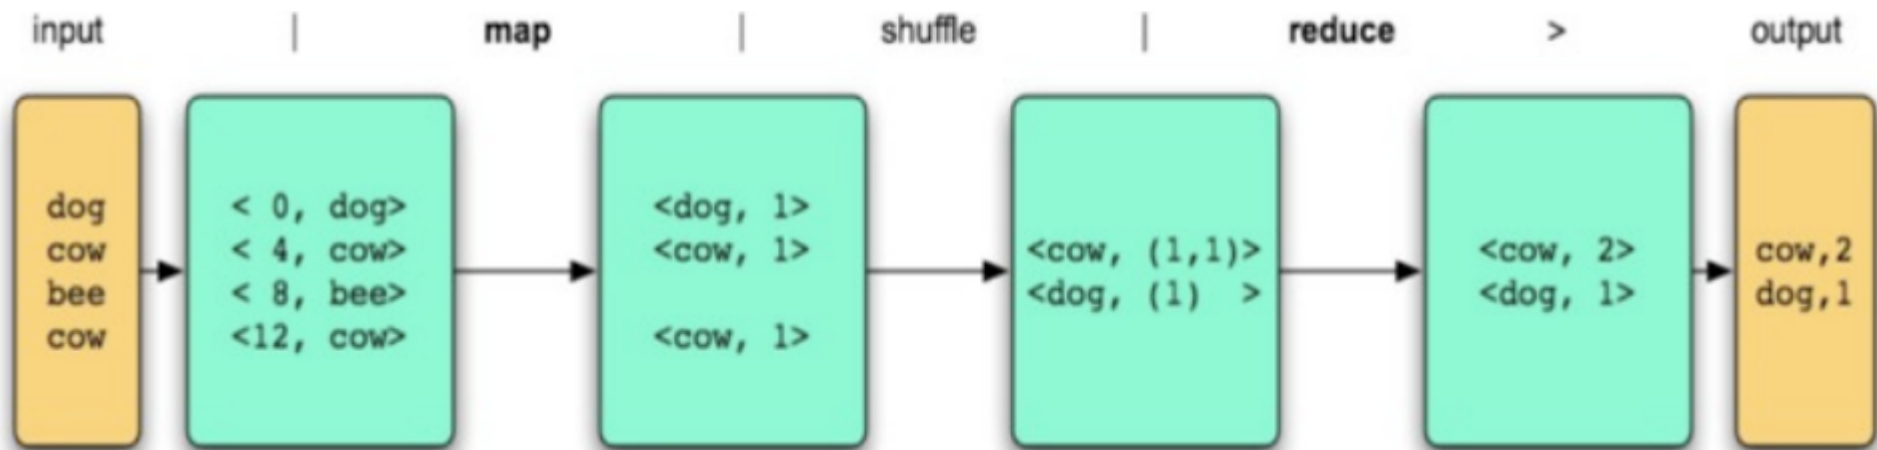
MapReduce的程式模型

- ▶ MapReduce函式示意

- Map (K1, V1) \rightarrow list(K2, V2)

- Reduce (K2, list(V2)) \rightarrow list(K3, V3)

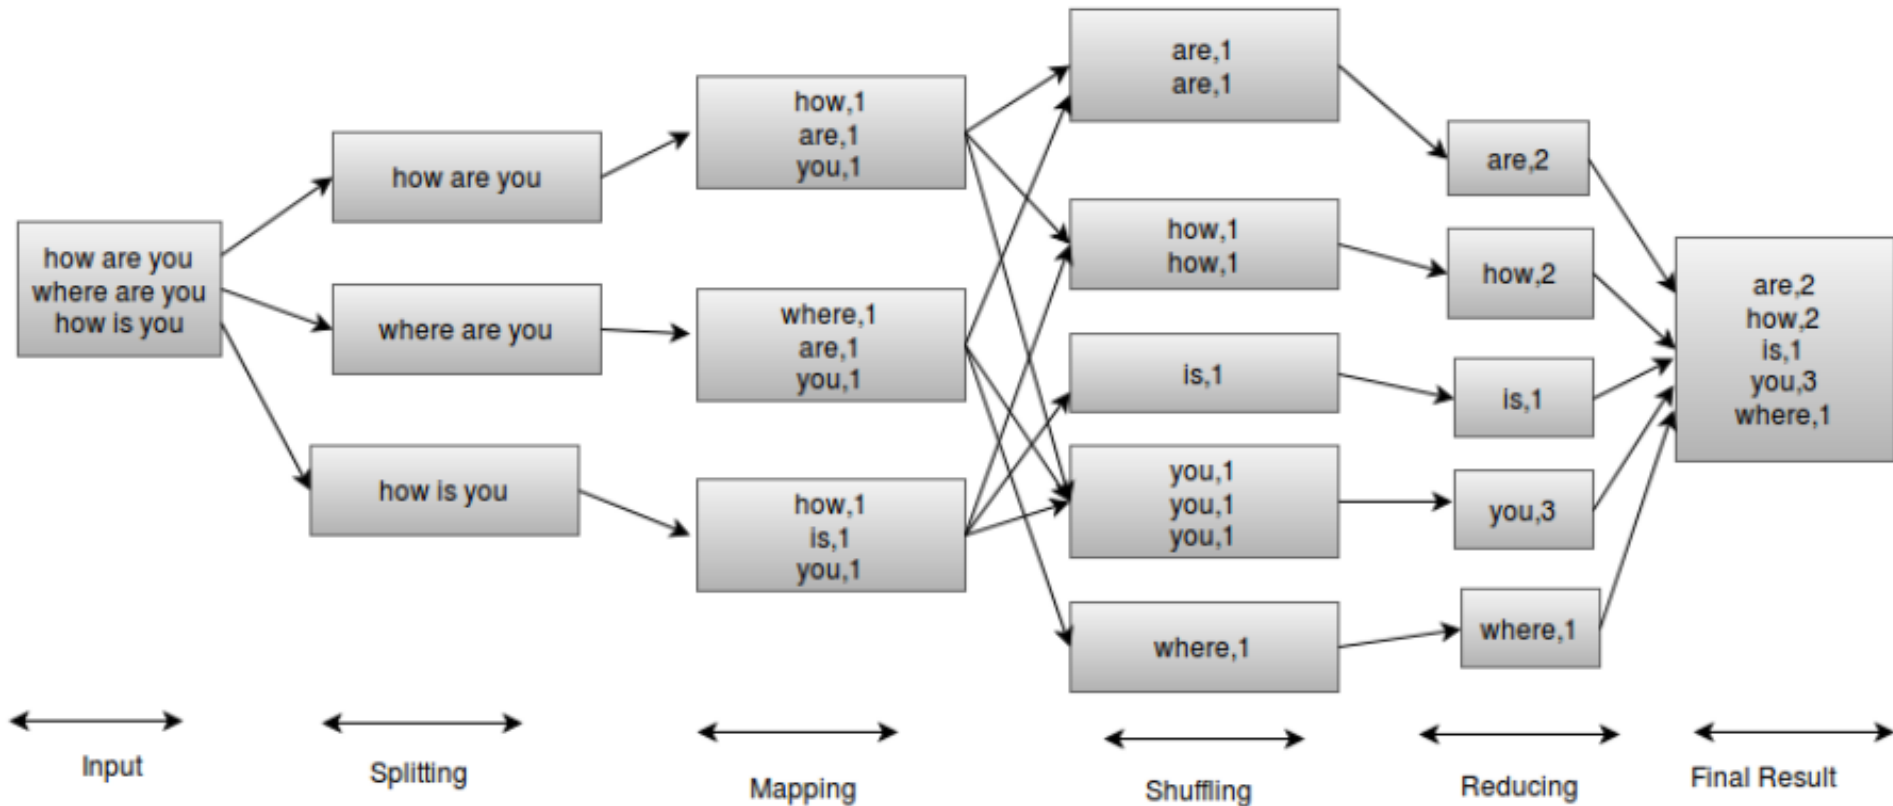
- ▶ 邏輯流程範例



MapReduce的程式模型(以WordCount為例)


- ▶ 試著如何在大量的文件集合中，計算每個字組(Word)的出現次數？
 - Map : (offset, line) -> (split) -> [word1, word2, word1] -> (collect) -> [(word1, 1),(word2, 1), (word1, 1)]
 - Shuffle(done by framework) : [(word1, 1),(word2, 1),(word1, 1)] -> (group) -> (group) -> [(word1, [1,1]),(word2, [1])]
 - Reduce : [(word1, [1,1]),(word2, [1])] -> (sum) -> [(word1, 2),(word2, 1)]

MapReduce的程式模型(以WordCount為例)



MapReduce In Hadoop

- ▶ Apache Hadoop實作Google的MapReduce
 - 提供Open Source的MapReduce框架(free for use!)
 - 以Java作為原生語言
 - 以HDFS作為輸出 / 輸入 / 中繼資料的儲存系統



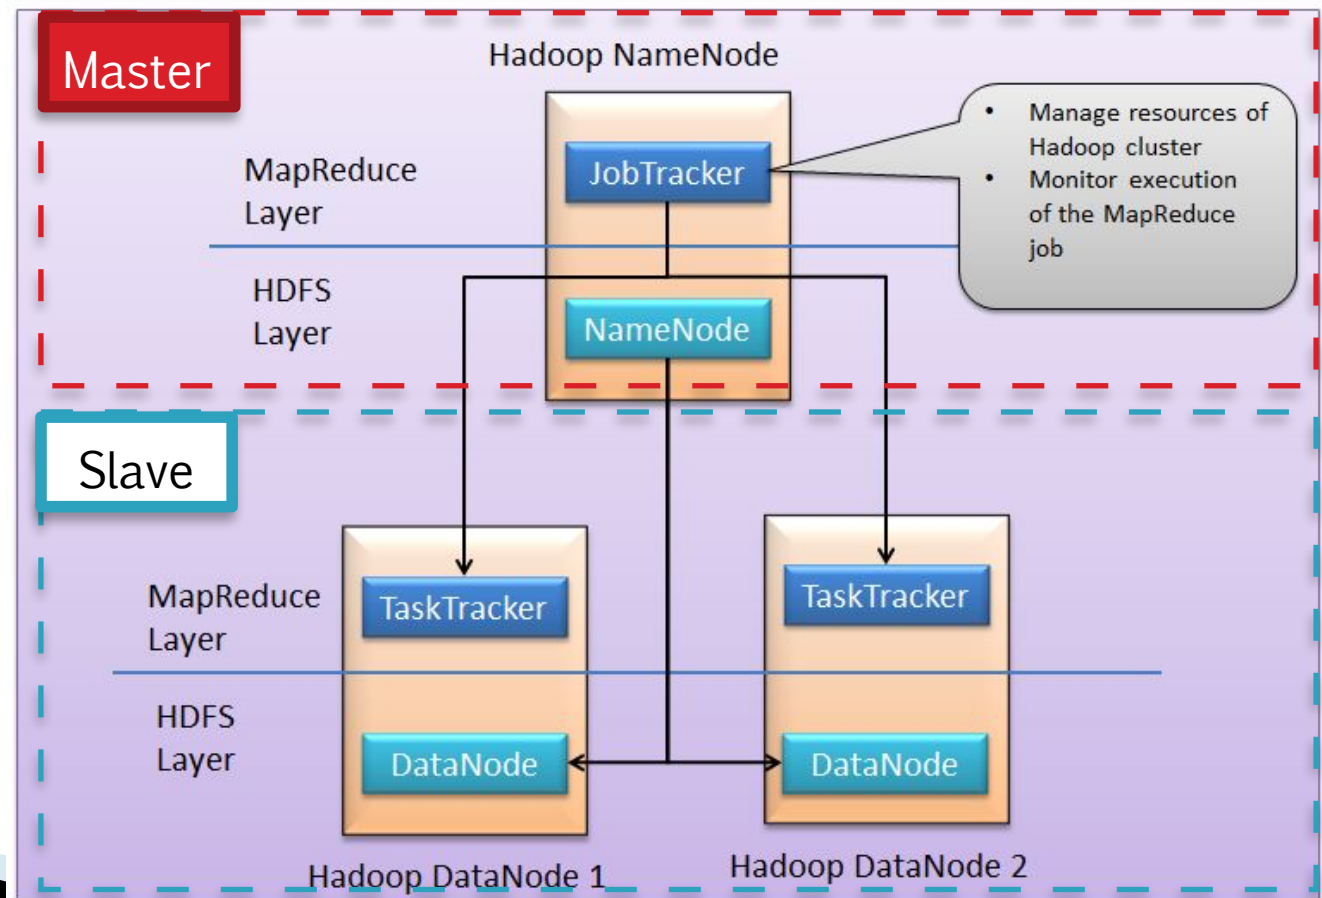
Hadoop MapReduce(V1)

▶ Master Daemon

- JobTracker
- 工作調配

▶ Slave Daemon

- TaskTracker
- 工作執行
- MapReduce執行



MapReduce V1 的缺點

▶ 延展性問題 / 效能問題

- JobTracker / NameNode只有一個，能夠管理的TaskTracker有限，縱使有大量的DataNode亦無用武之地

▶ 可及性問題

- NameNode有單點失敗問題

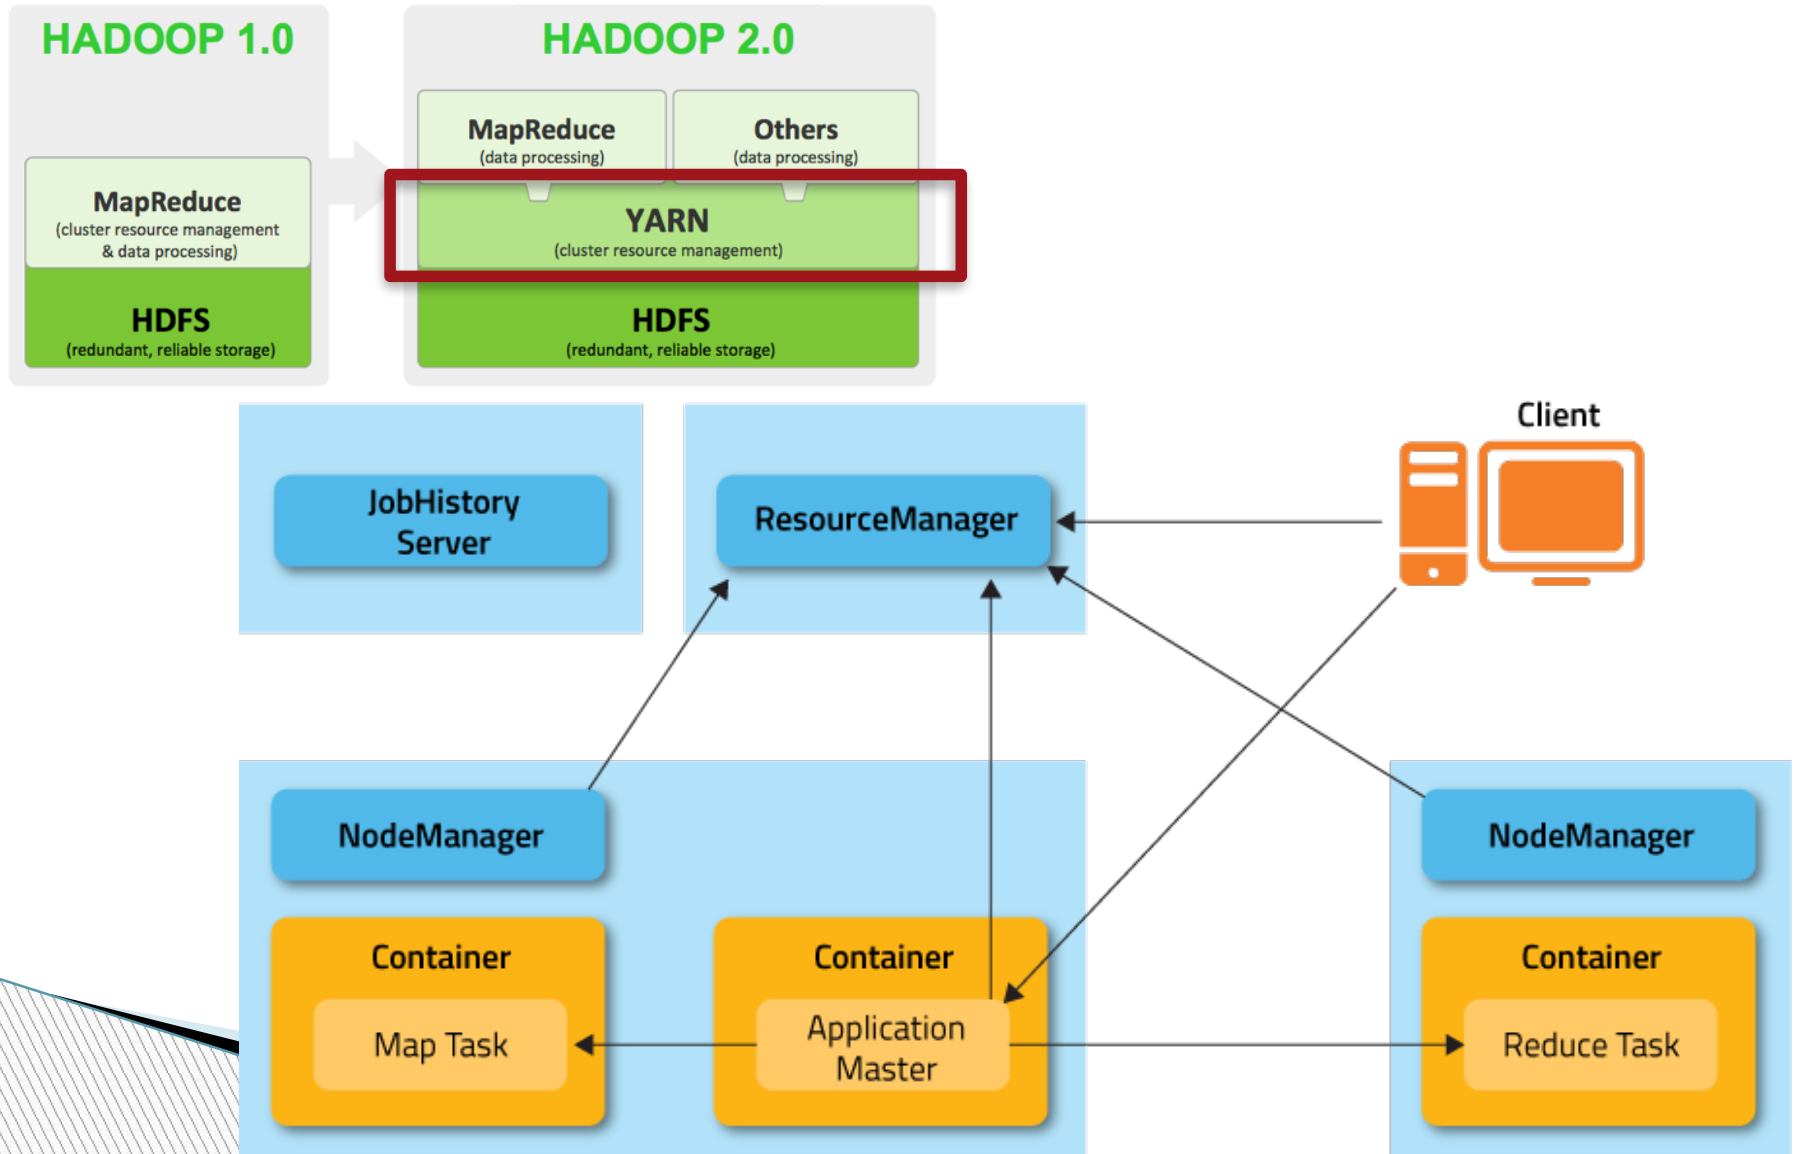
▶ 資源運用問題

- 在V1的設計中，Job的分配是考量Node的Task數進行分配，而非依據Node的CPU / Disk使用狀況，易有資源運用問題
- V1強制配置Map Slot及Reduce Slot，若Job只有Map Task會造Reduce Slot浪費

▶ 與異質系統的藕合問題

- V1的JobTracker僅支援MapReduce的應用程式，非MapReduce的框架(如Spark)無法運作

MapReduce V2 / YARN架構



MapReduce V2 / YARN設計

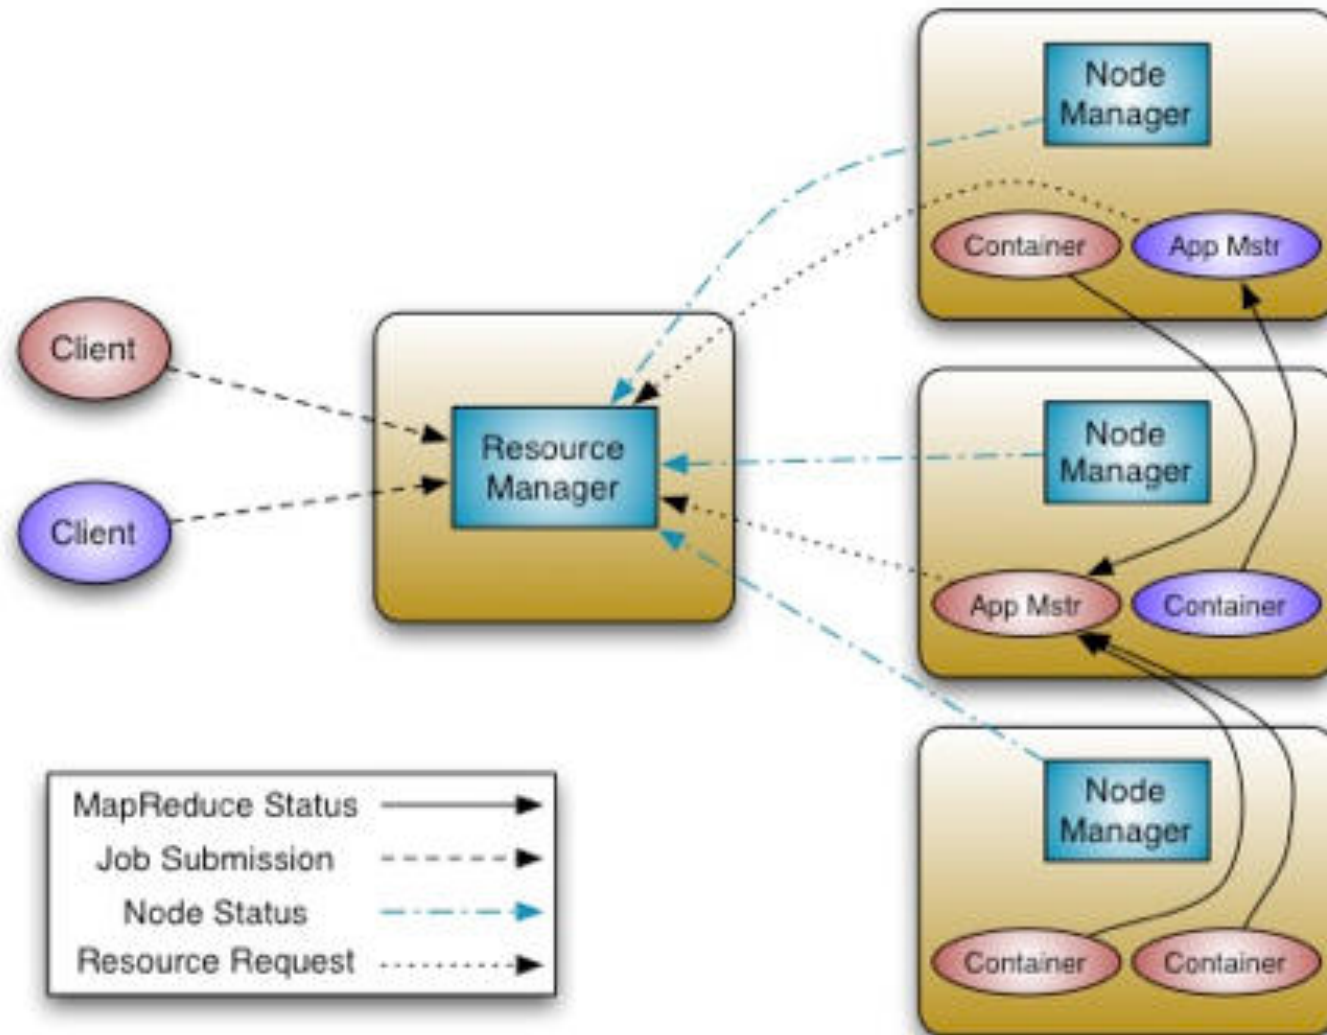
▶ 改善V1的問題

- 將Job Tracker的作業控制及資源管理分開
 - Resource Manager：資源管理與調度
 - Application Master：Task執行與控制
 - Node Manager：Slave上的資源與任務管理器
 - Container：依據需求所動態配置的資源
- 透過YARN支援異質運算框架

MapReduce V2 / YARN運作

- ▶ Resource Manager
 - 資源管理與調度(依Application Master的需求，指示Node Manager建立Container及Task)
- ▶ Application Master
 - 依運算特性向Resource Manager要求資源
 - 監控Task運作狀況
- ▶ Node Manager
 - 依Resource Manager要求配置資源並建立Container及Task


MapReduce V2 / YARN運作



Hadoop MapReduce程式開發

- ▶ 以WordCount為例
 - 開發工具 — Scala IDE (<http://scala-ide.org>)
 - 建立Java Project
 - 設定build Path，引入Hadoop相關library
 - Hadoop jar檔存放於 \$ HADOOP_HOME / share底下
 - common/hadoop-common.jar
 - common/lib/*.jar
 - hdfs/hadoop-hdfs.jar
 - mapreduce/hadoop-mapreduce-client-*.jar
 - yarn/hadoop-yarn-*.jar

WordCount執行過程示意



Hadoop MapReduce的型態

▶ Writable介面

- 在org.apache.hadoop.io套件下
 - IntWritable , LongWritable , Text
 - 對應到Java原生型態IntWritable => Int
- 主要功能
 - 資料序列化及反序列化，方便資料交換
- 所有Mapper與Reducer都必需**使用Writable Interface**作為輸入參數

Hadoop MapReduce的型態

▶ Input

- 在org.apache.hadoop.mapreduce.lib.input套件下
 - InputFormat - FileInputFormat, TextInputFormat
 - RecordReader - LineRecordReader
- 主要功能
 - 將資料在Map階段執行分割(Splits)，確定Map Task個數
 - 產生RecordReader，以從Splits產生一連串key / value
- InputFormat及RecordReader是Input及key / value的溝通橋樑

Hadoop MapReduce的型態

- ▶ Output

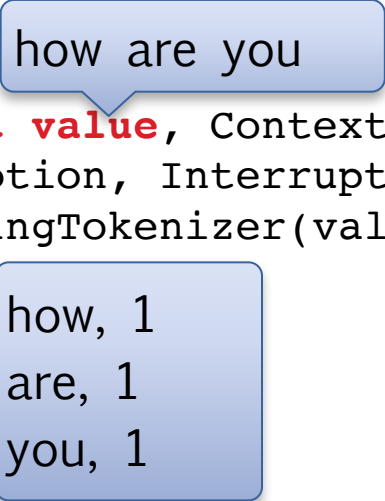
- 在org.apache.hadoop.mapreduce.lib.output套件下
 - OutputFormat - FileOutputFormat, TextOutputFormat
 - RecordWriter
- 主要功能
 - 將產的Key / value輸出至檔案系統中

WordCount程式分析 – Mapper

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```



WordCount程式分析－Reducer

```
public static class IntSumReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    private IntWritable result = new IntWritable();
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, are, 1 IOException
    {
        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result); are, 2
    }
}
```

對應Mapper的輸出

宣告Reducer的輸出

are, 1
are, 1

WordCount程式分析 – Driver

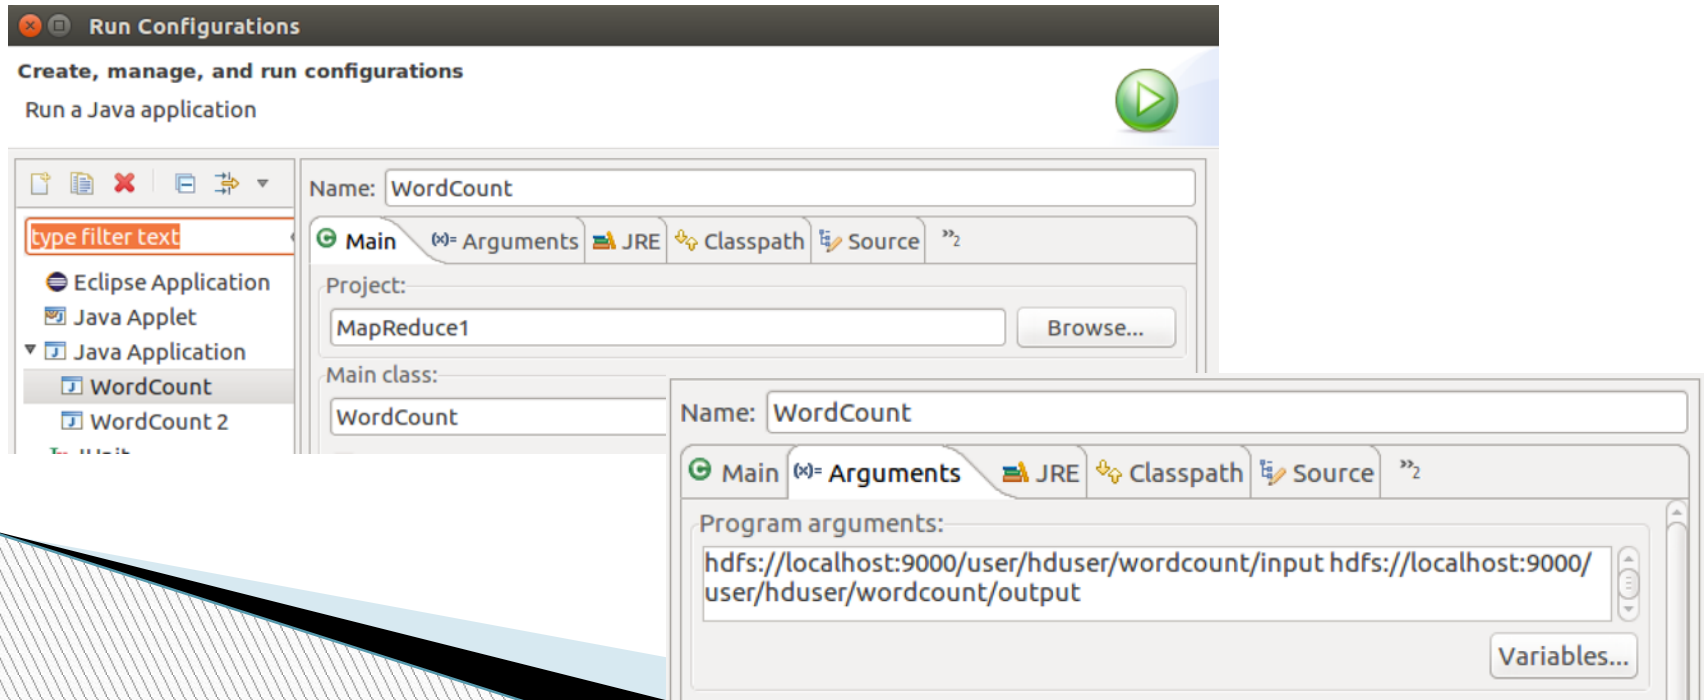
```
public static void main(String[] args) throws Exception {  
    Configuration conf = new Configuration();  
    Job job = Job.getInstance(conf, "word count");  
    job.setJarByClass(WordCount.class);  
    job.setMapperClass(TokenizerMapper.class);  
    job.setCombinerClass(IntSumReducer.class);  
    job.setReducerClass(IntSumReducer.class);  
    job.setOutputKeyClass(Text.class);  
    job.setOutputValueClass(IntWritable.class);  
    FileInputFormat.addInputPath(job, new Path(args[0]));  
    FileOutputFormat.setOutputPath(job, new Path(args[1]));  
    System.exit(job.waitForCompletion(true) ? 0 : 1);  
}
```

WordCount執行－前置作業

- ▶ 下載蓋茲堡宣言 (<https://github.com/yclee0418/hadoopTeach/blob/master/mapReduce/gettysburg.txt>)
- ▶ 在HDFS建立input資料夾
 - `hadoop fs -mkdir -p /user/hduser/wordcount/input`
- ▶ 上傳檔案至HDSF
 - `hadoop fs -copyFromLocal gettysburg.txt /user/hduser/wordcount/input`

WordCount執行 – IDE Debug / Run

- ▶ 設定debug / run configuration
- ▶ 設定Arguments
 - args[0] - hfs://localhost:9000/user/hduser/wordcount/input
 - args[1] - hfs://localhost:9000/user/hduser/wordcount/output



WordCount執行－Hadoop指令

- ▶ 在IDE中將Java Project匯出成Jar檔(EX : wc.jar)
- ▶ Hadoop指令
 - `hadoop jar wc.jar WordCount /user/hduser/wordcount/input /user/hduser/wordcount/output`

```
hduser@spark-single:~/Downloads$ hadoop jar wc.jar WordCount /user/hduser/wordcount/input /user/hduser/wordcount/output
16/12/07 21:42:26 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032
16/12/07 21:42:26 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed
. Implement the Tool interface and execute your application with ToolRunner to remedy this.
16/12/07 21:42:27 INFO input.FileInputFormat: Total input paths to process : 1
16/12/07 21:42:27 INFO mapreduce.JobSubmitter: number of splits:1
16/12/07 21:42:27 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1480943616948_0002
16/12/07 21:42:28 INFO impl.YarnClientImpl: Submitted application application_1480943616948_0002
16/12/07 21:42:28 INFO mapreduce.Job: The url to track the job: http://spark-single:8088/proxy/application_1480943616948_0002/
16/12/07 21:42:28 INFO mapreduce.Job: Running job: job_1480943616948_0002
16/12/07 21:42:38 INFO mapreduce.Job: Job job_1480943616948_0002 running in uber mode : false
16/12/07 21:42:38 INFO mapreduce.Job:  map 0% reduce 0%
16/12/07 21:42:44 INFO mapreduce.Job:  map 100% reduce 0%
16/12/07 21:42:50 INFO mapreduce.Job:  map 100% reduce 100%
16/12/07 21:42:51 INFO mapreduce.Job: Job job_1480943616948_0002 completed successfully
16/12/07 21:42:51 INFO mapreduce.Job: Counters: 49
```