

Outline

Module 1 : 大數據簡介

Module 2 : Hadoop Ecosystem介紹

Module 3 : Hadoop 平台安裝

Module 4 : Hadoop 分散式檔案系統 (HDFS)

Module 5 : Hadoop MapReduce

Module 6 : Apache Hive

Module 7 : Sqoop與Flume

Module 8 : Apache Spark

Module 9 : Spark 平台安裝

Module 10 : RDD — Resilient distributed dataset

Module 11 : Scala 程式開發基礎

Module 12 : Spark SQL 及 DataFrame

Module 13 : Spark 機器學習函式庫(MLlib)



Spark SQL 的由來

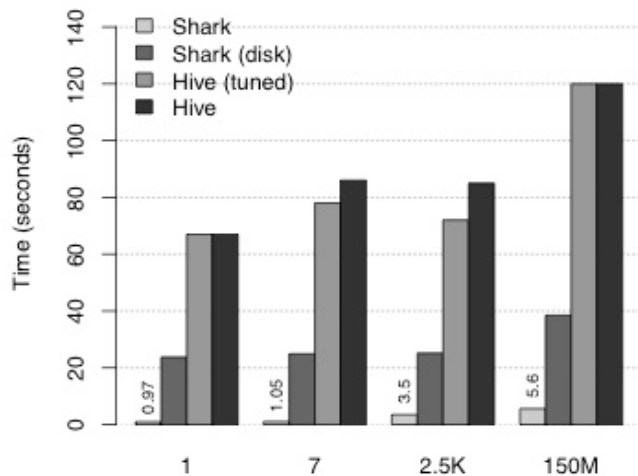
- ▶ Hive提供**熟悉SQL**但又不理解的MapReduce的技術人員快速上手的工具
- ▶ 但MapReduce計算過程中**大量的磁碟I / O造成效能瓶頸**，SQL-on-Hadoop的工具因應而生
 - MapR的Drill
 - Cloudera的Impala
 - **Shark(Spark SQL的前身)**
- ▶ Spark平台的SQL專案原為Shark，但Shark與Hive相依性太高；故另啟SparkSQL專案以取得更大彈性

Spark SQL 介紹

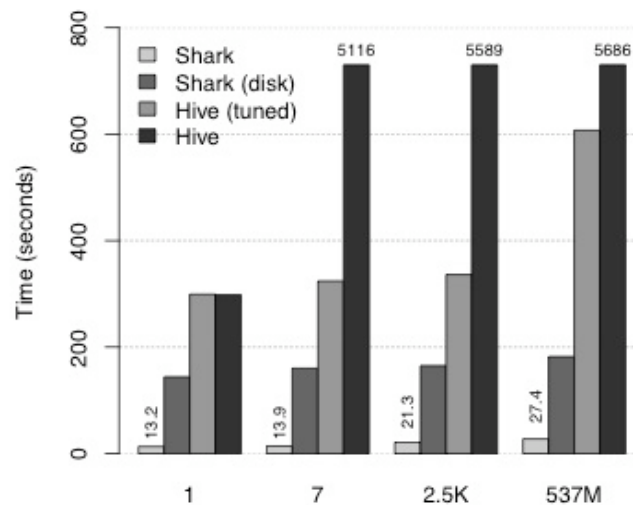
- ▶ Spark SQL is Apache Spark's module for working with **structured data**.
- ▶ **無縫與Spark程式整合**
- ▶ 提供**SQL-Like語法或DataFrame API**操作RDD
 - 降低學習門檻(相較Scala RDD API)
 - 提供更佳效能(In memory、SQL最佳化)

```
context = HiveContext(sc)
results = context.sql(
    "SELECT * FROM people")
names = results.map(p => p.name)
```

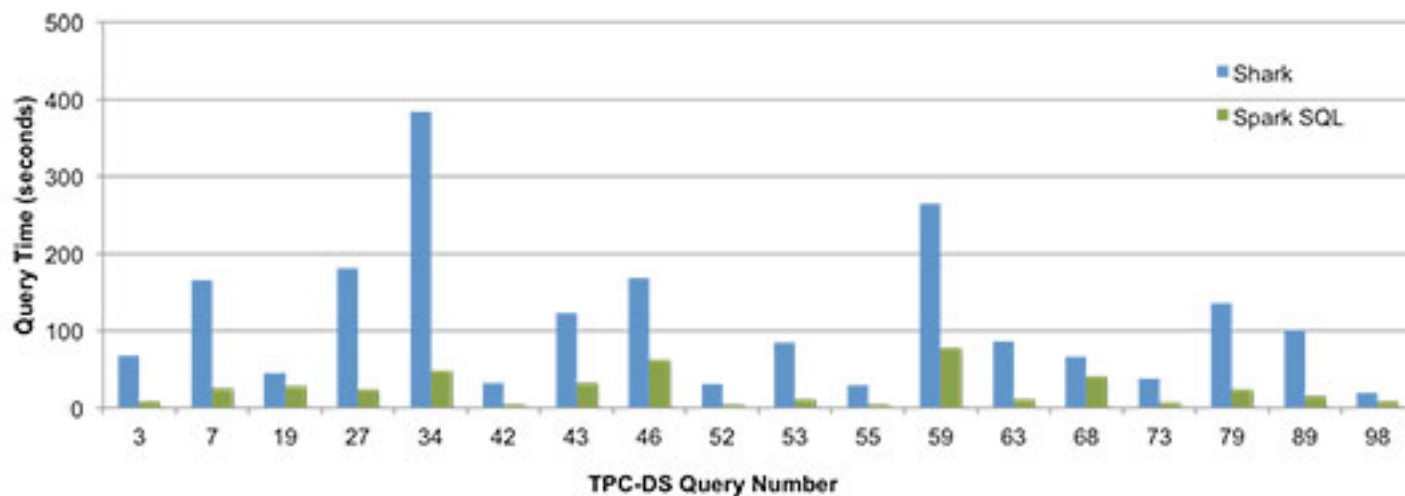
效能比較



TPC-H 100GB



TPC-H 1TB



Spark SQL 介紹

- ▶ 提供**單一程式介面存取多種資料來源**

- Hive、Parquet


- JSON、JDBC

- ▶ 與既有Hive應用無縫接軌

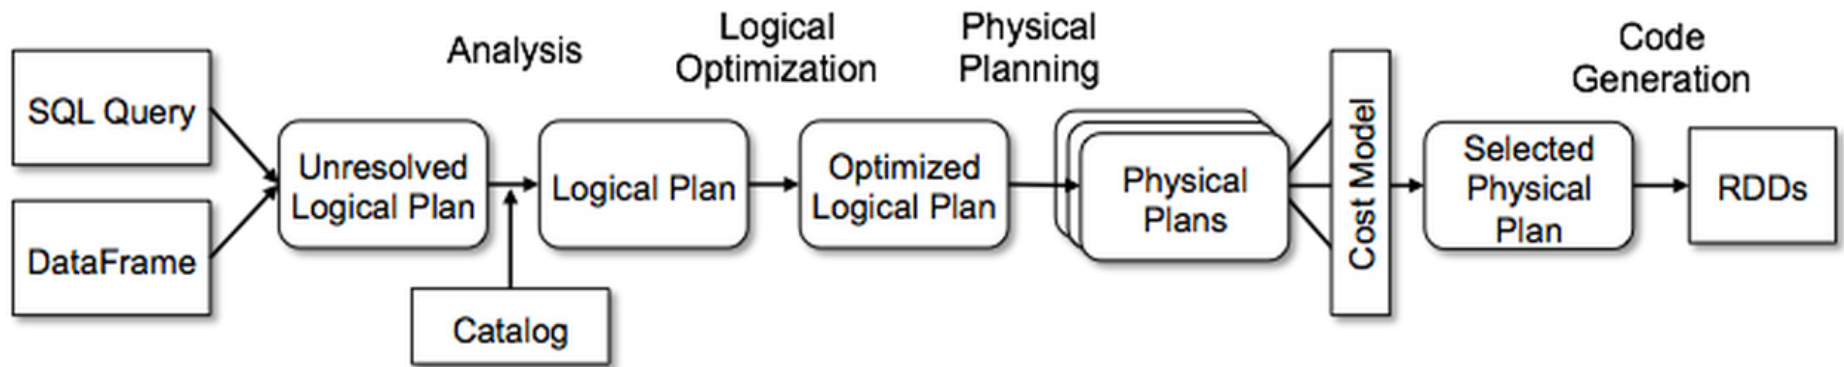
- 可存取Hive Table

- UDFs、SerDes直接使用

```
context.jsonFile("s3n://...")  
  .registerTempTable("json")  
results = context.sql(  
  ""SELECT * FROM people JOIN json ...""
```



Spark SQL Query Planning



Spark SQL支援的資料來源

Built-In

{ JSON }



External



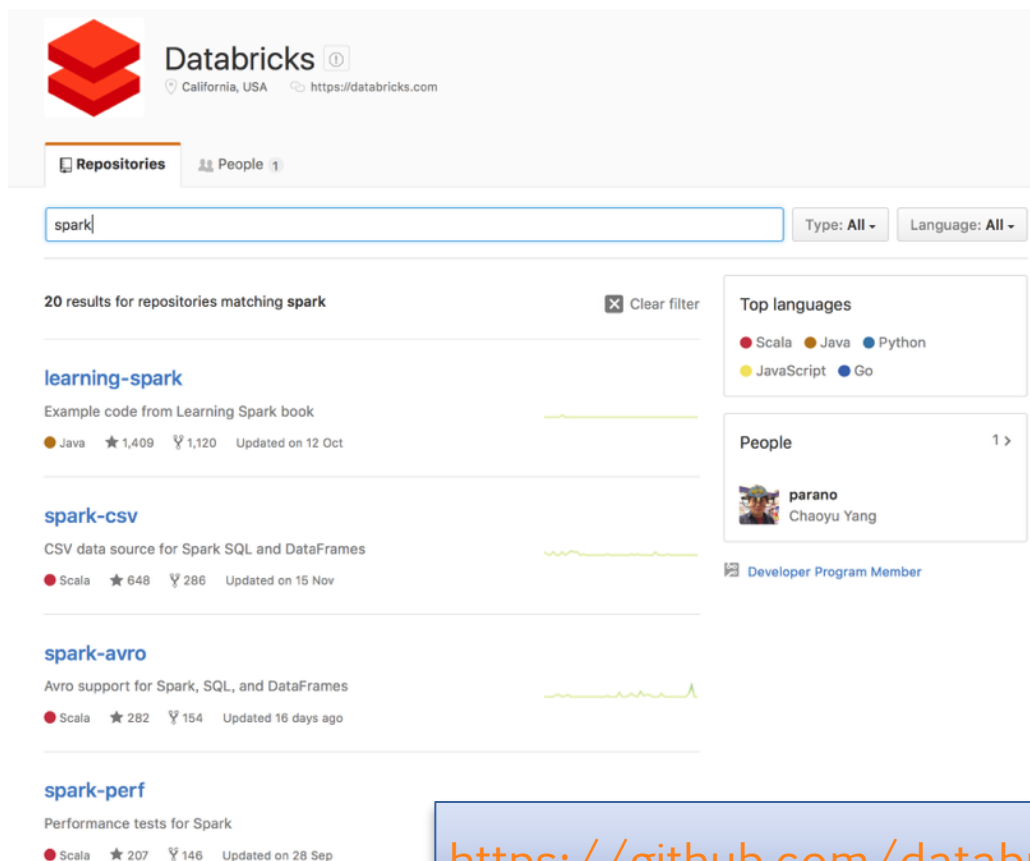
elasticsearch.



and more...

Spark SQL支援的資料來源

- ▶ External的資料來源可由Databricks的Github取得



The screenshot shows the Databricks GitHub profile page. The header includes the Databricks logo, name, location (California, USA), and website (https://databricks.com). Below the header, there are tabs for 'Repositories' and 'People'. A search bar contains the text 'spark'. The results show 20 repositories matching 'spark'. The first three repositories are listed: 'learning-spark' (Example code from Learning Spark book, Java, 1,409 stars, 1,120 forks, updated on 12 Oct), 'spark-csv' (CSV data source for Spark SQL and DataFrames, Scala, 648 stars, 286 forks, updated on 15 Nov), and 'spark-avro' (Avro support for Spark, SQL, and DataFrames, Scala, 282 stars, 154 forks, updated 16 days ago). The fourth repository is 'spark-perf' (Performance tests for Spark, Scala, 207 stars, 146 forks, updated on 28 Sep). On the right side, there are sections for 'Top languages' (Scala, Java, Python, JavaScript, Go) and 'People' (parano, Chaoyu Yang, Developer Program Member). At the bottom right, there is a blue box with the URL <https://github.com/databricks>.

SparkSQL in Spark-Shell

- ▶ SparkSQL已內建於Spark平台中，不需Hive即可直接使用
- ▶ 在程式中，SparkSQL皆透過**SparkSession**操作
- ▶ 在spark-shell啟動後，兩個重要的類別會自動被初始化
 - **sc變數**：SparkContext的Instance，為操作RDD的接口
 - **spark變數**：SparkSession的Instance，為操作DataFram API及SQL的接口

SparkSQL in Spark-Shell

直接啟動spark-shell

```
hduser@spark-single:~/Downloads$ spark-shell
Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel).
16/12/24 11:30:34 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
16/12/24 11:30:34 WARN Utils: Your hostname, spark-single resolves to a loopback address: 127.0.1.1; using 10.0.2.15 instead (on interface enp0s3)
16/12/24 11:30:34 WARN Utils: Set SPARK_LOCAL_IP if you need to bind to another address
16/12/24 11:30:35 WARN SparkContext: Use an existing SparkContext, some configuration may not take effect.
Spark context Web UI available at http://10.0.2.15:4040
Spark context available as 'sc' (hostname = spark-single, port = 7077).
Spark session available as 'spark'.
Welcome to
```

SparkContext被初始化於sc變數

SparkSession被初始化於spark變數

```

  ____  _
 / ___|| | | |
| |___| |_| |
 \___ \|  __/
      | |
      |_|

version 2.0.2
```

```
Using Scala version 2.11.8 (OpenJDK 64-Bit Server VM, Java 1.8.0_111)
Type in expressions to have them evaluated.
Type :help for more information.

scala> █
```

SparkSession

- ▶ SparkSession作為SparkSQL的程式接口(以下簡稱為spark)，具有資料讀取、DataFrame操作及SQL指令操作功能
- ▶ spark.catalog
 - 新增及刪除SparkSQL中Table及DataBase的介面
 - ex：spark.catalog.listDatabases.show

```
scala> spark.catalog.  
cacheTable          dropTempView        listFunctions        setCurrentDatabase  
clearCache          isCached            listTables           uncacheTable  
createExternalTable listColumns          refreshByPath  
currentDatabase     listDatabases        refreshTable
```

Ref: <http://spark.apache.org/docs/latest/api/scala/index.html#org.apache.spark.sql.SparkSession>

SparkSession

- ▶ spark.read
 - 讀取外部資料來源成為DataFrame
 - ex : `spark.read.csv("file:/home/hduser/ratings.txt")`
 - 內建支援的資料來源
 - csv、jdbc、parquet、text file、json
 - 以option設定讀取參數(如csv的欄位分割符號)
 - ex : `spark.read.option("delimiter", "\t").csv(path)`

```
scala> spark.read.  
csv      jdbc      load      options    parquet    table    textFile  
format   json    option   orc        schema     text
```

SparkSession

- ▶ spark.sql
 - 以SQL指令操作SparkSQL的入口
 - ex : spark.sql("show databases").show
 - 在spark-shell中可省略spark.簡寫為sql("...")
 - SparkSQL的sql語法為SQL-92的子集
 - 支援project、where、order、join、group、having等語法

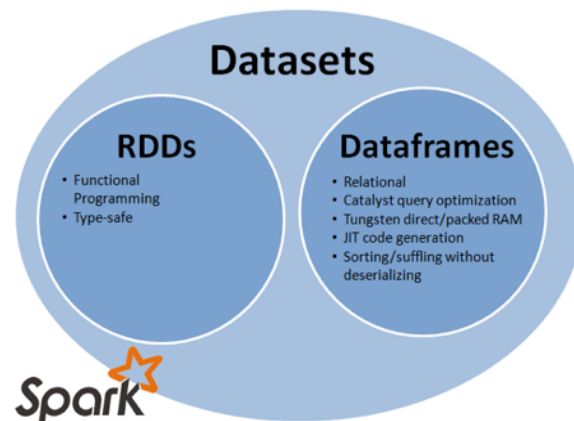
```
SELECT [DISTINCT] [column names]|[wildcard]
FROM [keyspace name.]table name
[JOIN clause table name ON join condition]
[WHERE condition]
[GROUP BY column name]
[HAVING conditions]
[ORDER BY column names [ASC | DSC]]
```

SparkSQL with Hive in Spark-Shell

- ▶ SparkSQL支援讀取及寫入Hive的Database及Table
- ▶ Spark 2.0.2的SparkSQL最高支援Hive 1.2.1版本
 - 使用更高的Hive版本會產生不相容錯誤
- ▶ 要在Spark-shell中存取Hive需額外進行以下步驟
 - 將以下3個檔案copy至/usr/local/Spark/conf下
 - hive-site.xml (位於/usr/local/hive/conf)
 - core-site.xml (位於/usr/local/hadoop/etc/hadoop)
 - hdfs-site.xml (位於/usr/local/hadoop/etc/hadoop)
 - 啟動spark-shell時要引入mysql connector的jar檔
 - spark-shell --jars /usr/local/hive/lib/mysql-connector-java-5.1.40-bin.jar

DataSet & DataFrame

- ▶ DataSet在SparkSQL中是裝載data的類別
 - 由RDD衍生而來，承襲RDD的優點
 - strong typing、支援lambda操作
 - 可與RDD相互轉換
 - 支援RDD的transformation操作(map, flatMap, filter, etc.)
- ▶ DataFrame是具有欄位名稱的DataSet
- ▶ DataFrame將是Spark MLlib未來唯一支援的輸入型態



DataTypes in SparkSQL

```
import org.apache.spark.sql.types._
```

Find full example code at "examples/src/main/scala/org/apache/spark/examples/sql/SparkSQLExample.scala" in the Spark repo.

Data type	Value type in Scala	API to access or create a data type
ByteType	Byte	ByteType
ShortType	Short	ShortType
IntegerType	Int	IntegerType
LongType	Long	LongType
FloatType	Float	FloatType
DoubleType	Double	DoubleType
DecimalType	java.math.BigDecimal	DecimalType
StringType	String	StringType
BinaryType	Array[Byte]	BinaryType
BooleanType	Boolean	BooleanType
TimestampType	java.sql.Timestamp	TimestampType
DateType	java.sql.Date	DateType
ArrayType	scala.collection.Seq	ArrayType(elementType, [containsNull]) Note: The default value of <i>containsNull</i> is <i>true</i> .

<http://spark.apache.org/docs/latest/sql-programming-guide.html#data-types>

Input & Output

- ▶ 透過sparkSession.read系列的方法來讀取不同資料來源
 - Local File System：讀取路徑加上file:的前置詞
 - ex：file:/usr/local/spark/examples/src/main/resources/people.json
 - HDFS：讀取路徑不加file:的前置詞
 - ex：/user/data/people.json
- ▶ 透過sparkSession.write系列的方法來輸出至不同目的地
 - write.save方法預設輸出為Parquet檔案
 - Local File System：讀取路徑加上file:的前置詞
 - HDFS：讀取路徑不加file:的前置詞

使用者自訂函式(UDF)

- ▶ 當DataFrame及SQL功能不敷使用時，可透過定義UDF來進行功能擴充
 - 以sparkSession.udf.register註冊UDF
 - 以lambda語法撰寫UDF
 - 直接在DataFrame API或SQL描述中使用

```
spark.udf.register("ageType", (age: Int)=> {  
    var aType = "unknown"  
    if (age > 0 && age < 20) aType = "teen"  
    if (age >=20 && age < 40) aType = "adult"  
    aType })
```

```
sql("select name, ageType(age) as ageType from people").show
```

從作中學：JSON & DataFrame API

- ▶ 在spark-shell中輸入以下指令
 - `val df = spark.read.json("file:/usr/local/spark/examples/src/main/resources/people.json")`
 - 觀察shell中顯示df包含的欄位及型態
 - `df.show()`
 - 查看people.json的內容

```
scala> val df = spark.read.json("file:/usr/local/spark/examples/src/main/resources/people.json")
df: org.apache.spark.sql.DataFrame = [age: bigint, name: string]

scala> df.show
+-----+-----+
| age | name |
+-----+-----+
| null | Michael |
| 30 | Andy |
| 19 | Justin |
+-----+-----+
```

從作中學：DataFram API

- ▶ `df.select($"name").show`
 - 只取name欄位
- ▶ `df.select($"name", $"age" + 1).show()`
 - 對欄位作變化($\text{age}=\text{age}+1$)
- ▶ `df.select($"name").filter($"age" > 10).show()`
 - 篩選年齡大於10歲的人名
- ▶ `df.groupBy($"age").count().show()`
 - groupBy及count的使用
- ▶ `df.groupBy($"age" > 10).count().show()`
 - groupBy的進階應用

Ref: <http://spark.apache.org/docs/latest/sql-programming-guide.html#datasets-and-dataframes>

從作中學：SQL語法

- ▶ `df.createOrReplaceTempView("people")`
 - 將DataFrame註冊為view
- ▶ `sql("select name from people").show`
 - 只取name欄位
- ▶ `sql("select name, age+1 from people").show`
 - 對欄位作變化($\text{age}=\text{age}+1$)
- ▶ `sql("select name from people where age > 10").show`
 - 篩選年齡大於10歲的人名
- ▶ `sql("select age,count(*) from people group by age").show`
 - groupBy及count的使用
- ▶ `sql("select age > 10,count(*) from people group by age>10").show`
 - groupBy及count的進階應用

從作中學：UDF

- ▶ 定義ageType：傳入age，0~20歲為teenager、20~40為adult、其餘為unknown
- ▶ 定義toUpper：傳入name，將name轉為大寫
- ▶ sql("select toUpper(name) as name, ageType(age) as ageType) from people")

```
spark.udf.register("ageType", (age: Int)=> {  
    var aType = "unknown"  
    if (age > 0 && age < 20) aType = "teen"  
    if (age >=20 && age < 40) aType = "adult"  
    aType })  
spark.udf.register("toUpper", (name: String)=> name.toUpperCase)  
sql("select toUpper(name), ageType(age) as ageType from people").show
```

實作演練：不同來源資料的JOIN

- ▶ 在Hive課程中我們已在Hive環境中建立Yelp的items資料表，接下來我們透過SparkSQL讀取ratings.txt檔案，並在SparkSQL中建立ratings View
- ▶ 利用SQL語法篩選出userid=0的使用者且評價 ≥ 4 的itemid及category

userid	itemid	ratings	itemid	category
0	0	4	0	Bowling Skating R...
0	1	5	1	Recreation Center...
0	7496	5	7496	Nail Salons Skin ...
0	7497	5	7497	Hair Salons Makeu...
0	7498	4	7498	Day Spas Skin Care
0	23691	5	23691	Mexican Bars
0	23694	4	23694	Lounges American ...
0	23695	4	23695	Sushi Bars Japanese

實作演練：不同來源資料的JOIN

- ▶ 啟動spark-shell
 - 注意要完成spark與hive的連結設定，並載入mysql-connector的jar
- ▶ 輸入以下指令
 - `val path="file:/home/hduser/Downloads/ratings.txt"`
 - `val ratingDF1=spark.read.option("header","true").option("inferSchema","true").option("delimiter",":").csv(path) #讀入ratings.txt為DF`
 - `ratingDF1.select("userid","itemid","ratings").createOrReplaceTempView("ratings") #在SparkSQL中建立VIEW`
 - `sql("show tables").show #顯示目前SparkSQL中的資料表`
 - `val userDf=sql("select userid,i.itemid,ratings,category from ratings r join items i on r.itemid=i.itemid where userid=0 and r.ratings>=4")`
 - 以itemid為key作join，篩選出userid=0的使用者且評價>=4的資料
 - `userDf.show`

實作演練：DataFrame內容的輸出

▶ 接續上頁實作

- `userDf.write.save("file:/home/hduser/Downloads/userDf.parquet")` #存成parquet檔案
- `userDf.write.csv("file:/home/hduser/Downloads/userDf.csv")` #存成csv檔案
- `userDf.write.json("file:/home/hduser/Downloads/userDf.json")` #存成json檔案
- `userDf.write.save("/user/userDf.parquet")` #存成parquet檔案，並寫入HDSF

實作演練：UDF

- ▶ 接續上頁實作
- ▶ 定義favorType：傳入ratings，5分為favorite、4分為not bad、其餘為soso
- ▶ sql("select *, favorType(ratings) from ratings where userid=0").show

```
spark.udf.register("favorType", (ratings: Int)=> {  
    var fType = "soso"  
    if (ratings == 5) fType = "favorite"  
    if (ratings == 4) fType = "not bad"  
    fType })
```

userid	itemid	ratings	UDF(ratings)
0	0	4	not bad
0	1	5	favorite
0	7495	3	soso
0	7496	5	favorite
0	7497	5	favorite
0	7498	4	not bad